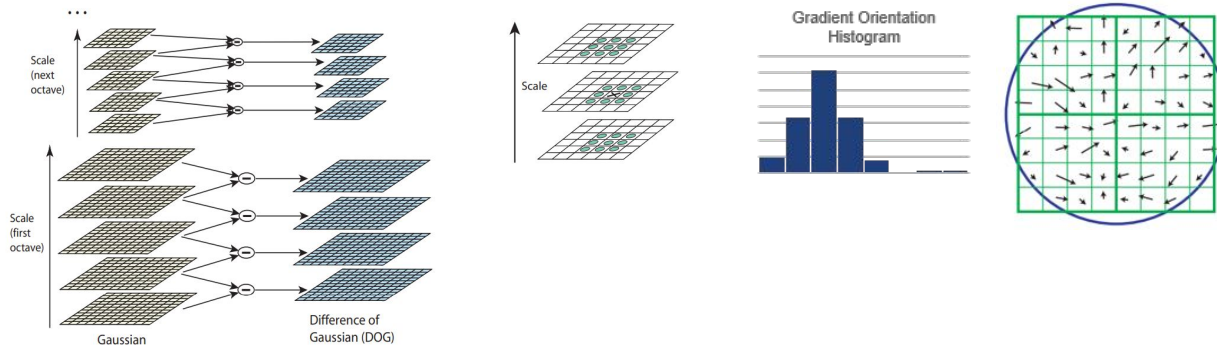
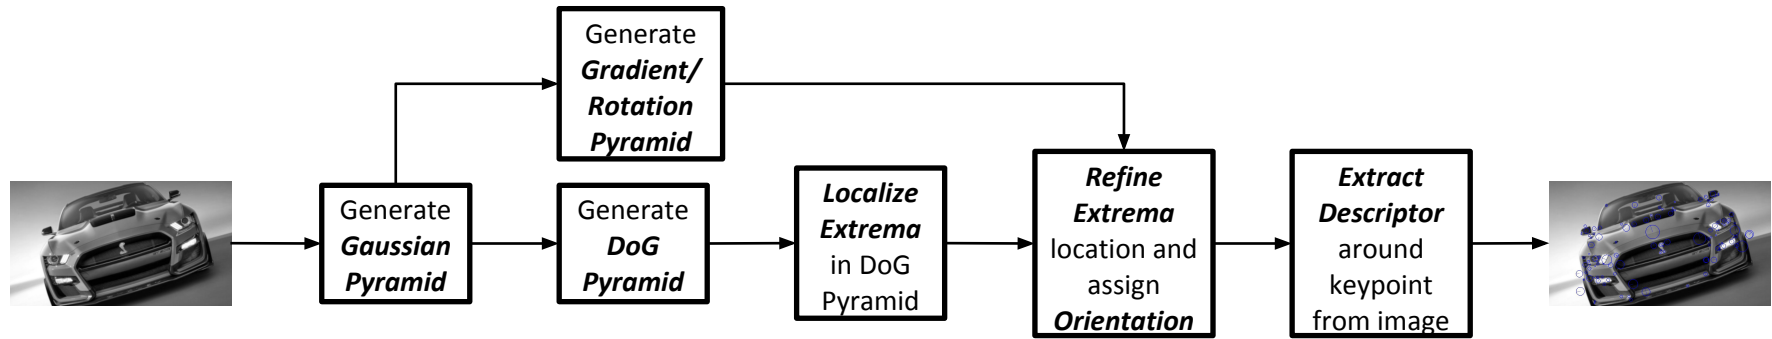


A SIFT Descriptor for Feature Matching

Advanced Systems Lab


Nicolas Hafner, Costanza Maria Improta, Zsombor Kalotay, Jan Leutwyler

Our Algorithm



Cost Analysis and Profiling

- **Flop count** and **memory accesses**



- **Bottlenecks:**


- Gaussian Pyramid
- Gradient and Rotation Pyramids

- **Reference Implementation:**

- ezSIFT

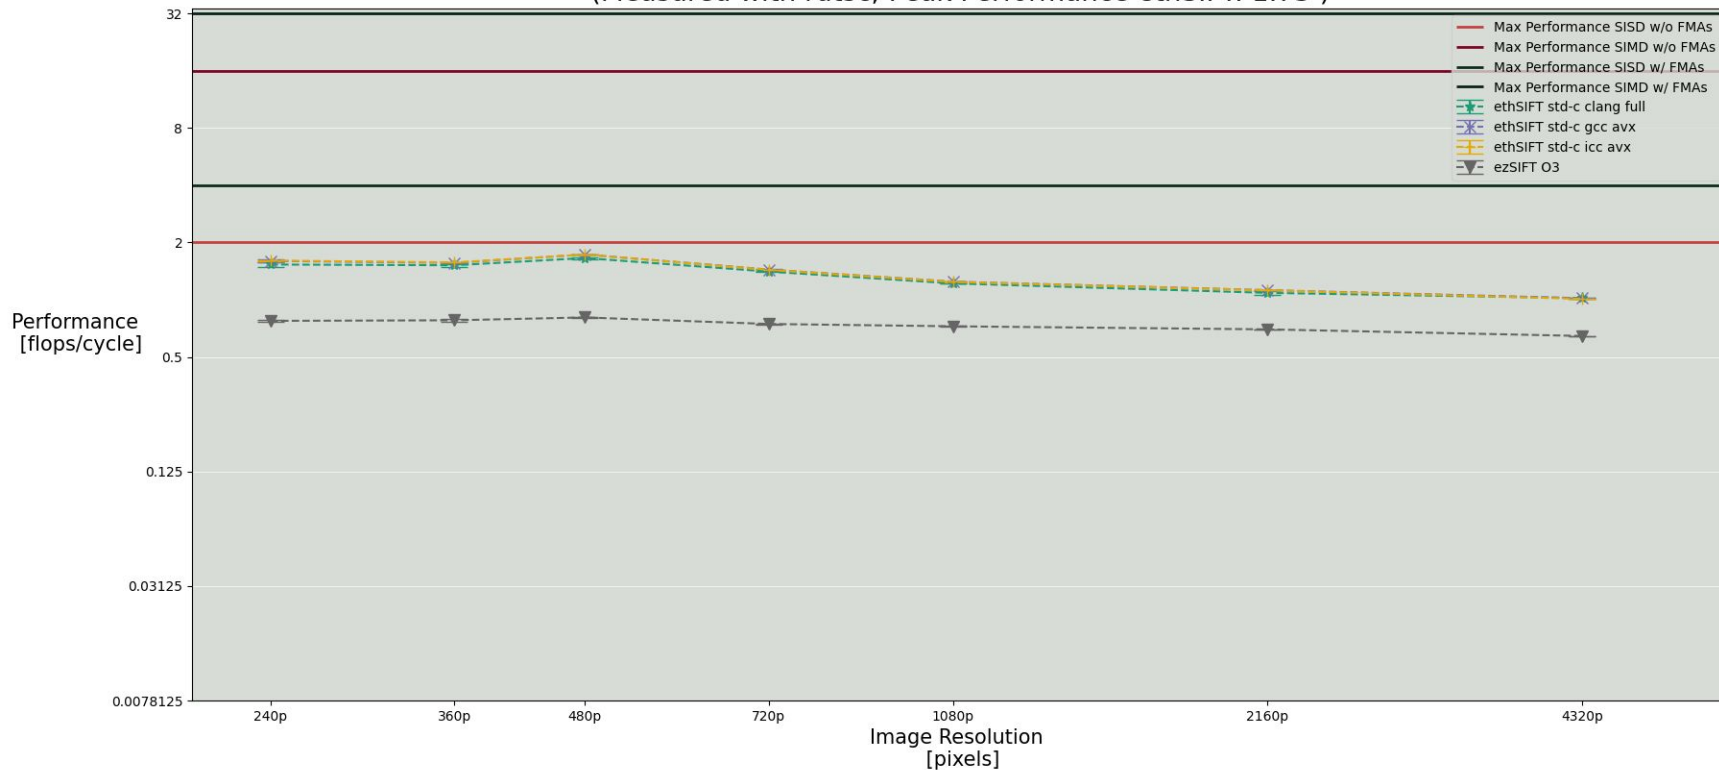
Baseline Implementation (Intel i7-8700K @ 4.4GHz)

Performance MeasureFull
(Measured with rdtsc; Peak Performance ethSIFT: 1.02)




Standard C optimisations (Intel i7-8700K @ 4.4GHz)

Performance MeasureFull
(Measured with rdtsc; Peak Performance ethSIFT: 1.73)




AVX optimisations (Intel i7-8700K @ 4.4GHz)

Performance MeasureFull
(Measured with rdtsc; Peak Performance ethSIFT: 3.24)




Some interesting plots (Intel i7-8700K @ 4.4GHz with ICC)

Performance Convolution
(Measured with rdtsc; Peak Performance ethSIFT: 6.06)




Some interesting plots (Intel i7-8700K @ 4.4GHz with ICC)

Performance GaussianPyramid
(Measured with rdtsc; Peak Performance ethSIFT: 6.49)




Some interesting plots (Intel i7-8700K @ 4.4GHz with ICC)

Performance GradientAndRotationPyramids
(Measured with rdtsc; Peak Performance ethSIFT: 5.16)



Some interesting plots (Intel i7-8700K @ 4.4GHz with ICC)

Performance ExtremaRefinement
(Measured with rdtsc; Peak Performance ethSIFT: 1.28)



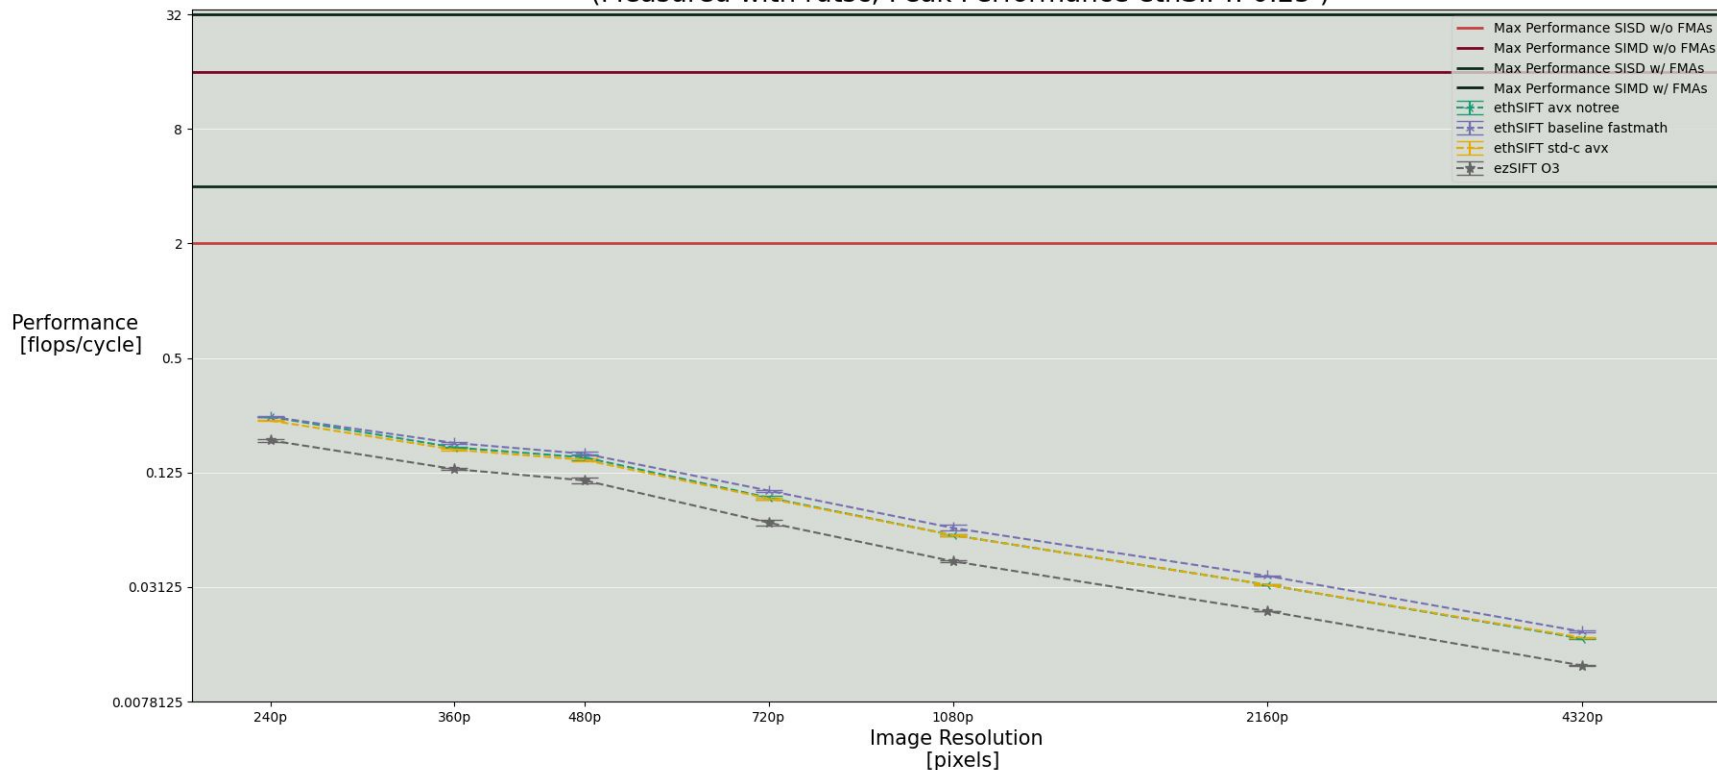
Some interesting plots (Intel i7-8700K @ 4.4GHz with ICC)

Performance Histogram
(Measured with rdtsc; Peak Performance ethSIFT: 1.06)



Some interesting plots (Intel i7-8700K @ 4.4GHz with ICC)

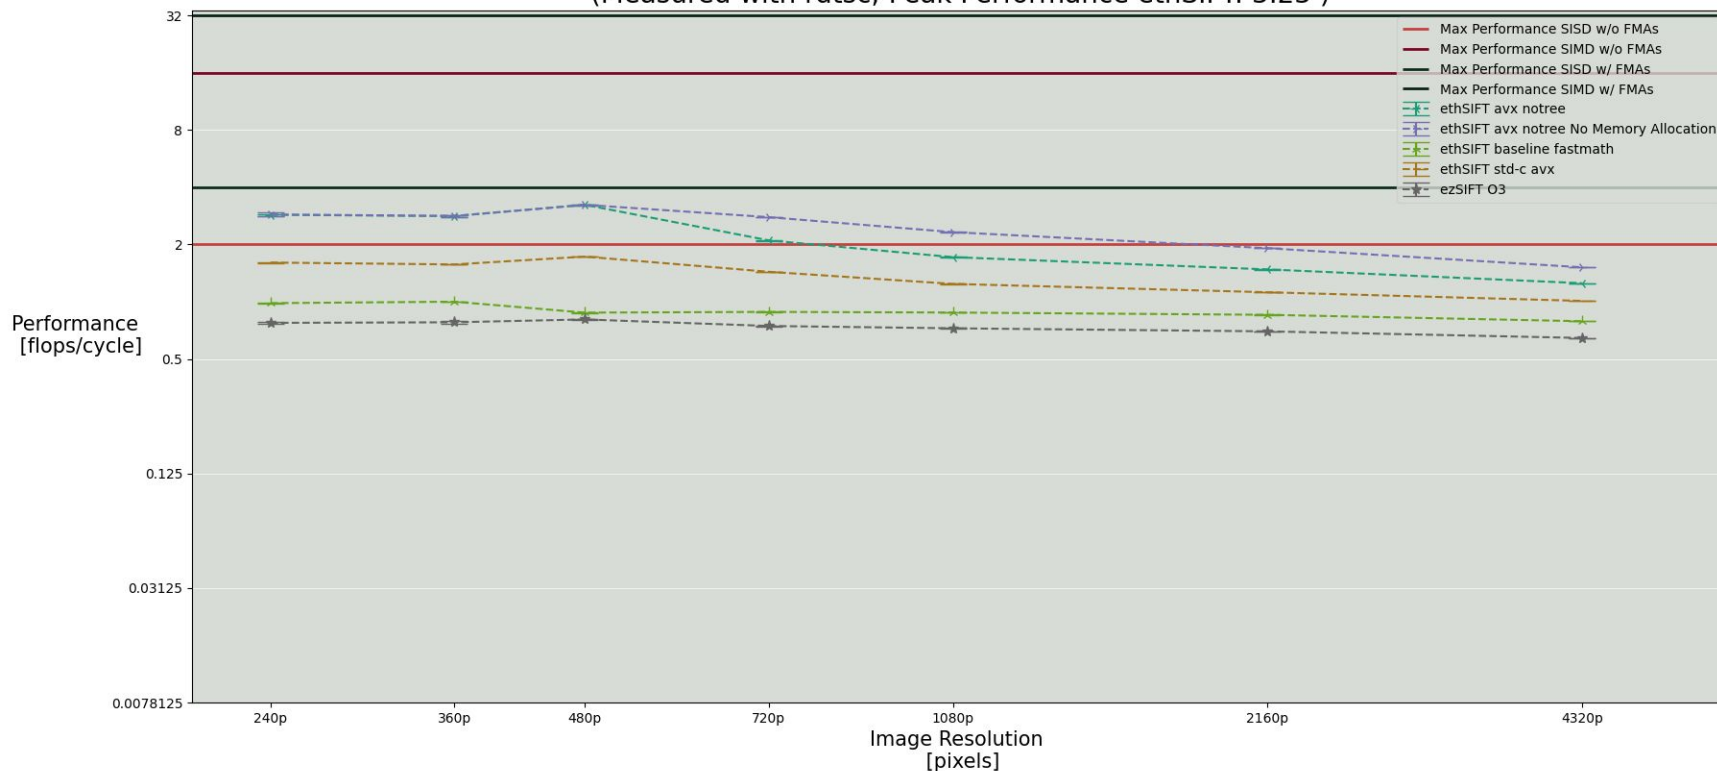
Performance KeypointDetection
(Measured with rdtsc; Peak Performance ethSIFT: 0.25)



Experimental Results (Intel i7-8700K @ 4.4GHz with ICC)

Overall Speedup:
~3x

Performance MeasureFull
(Measured with rdtsc; Peak Performance ethSIFT: 3.23)



Interesting Findings

★ **ezSIFT leaks** memory by default (manually fixed)

★ **Helped** performance:

- Variable **declarations close to point of use**
- Declaring **const**

★ **Worsened** performance:

- **size_t** instead of int
- Forcing **memory** to be **paged** when allocated

★ **Sometimes** helped sometimes hurt:

- **Prefetching** memory
- Compiler **flags**

Without optimisations

- 1) GCC
- 2) Clang
- 3) ICC

With optimisations

- 1) ICC
- 2) Clang
- 3) GCC

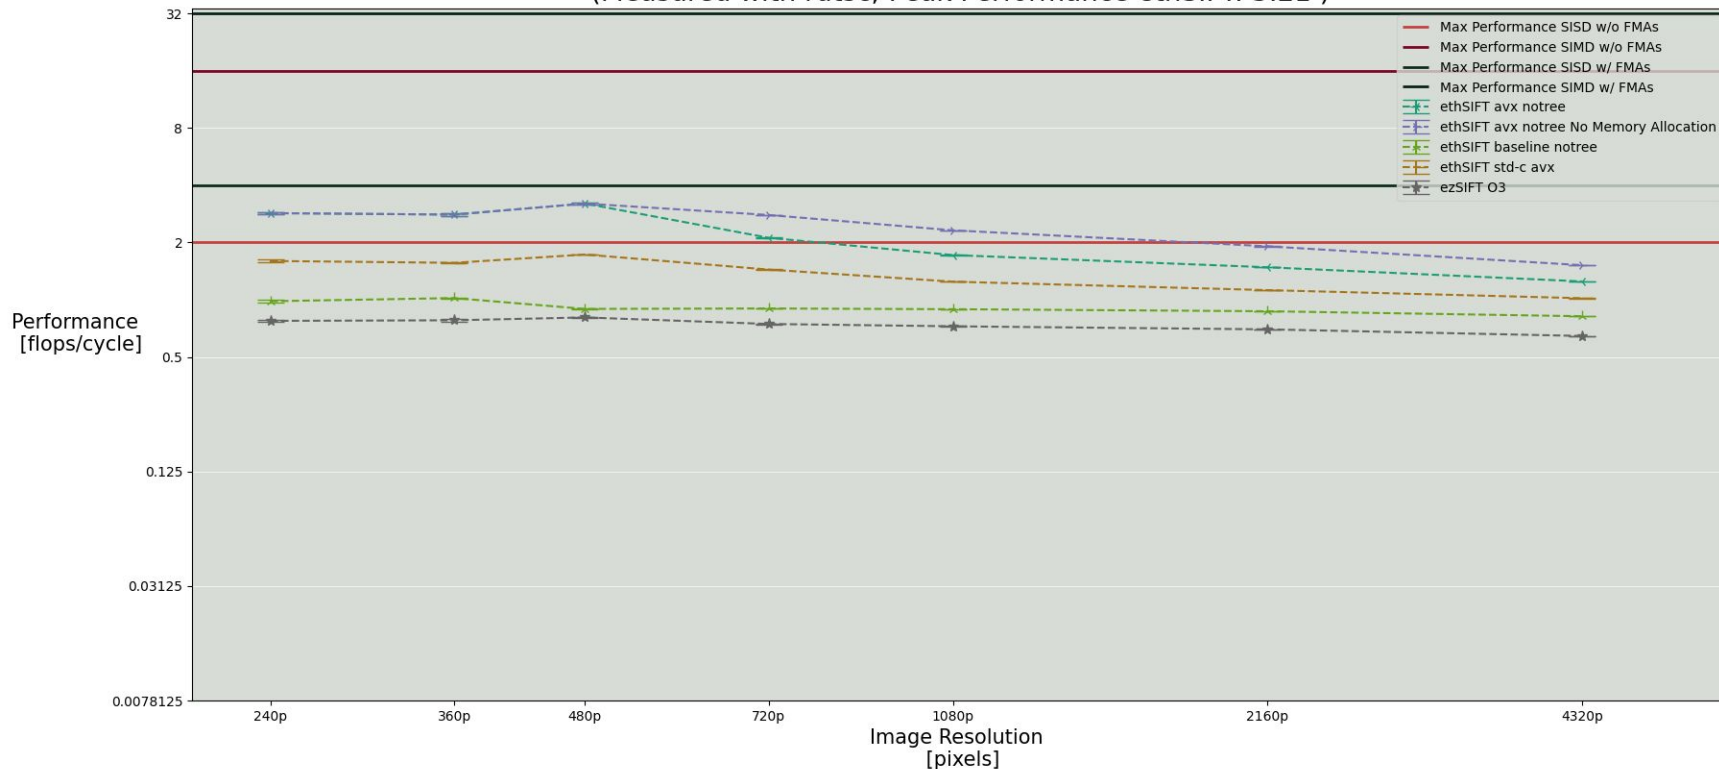
In Conclusion

- **SIFT is difficult to optimise**
- **Gaussian** kernel convolution quickly becomes the **biggest bottleneck**
- In parts achieved **up to 10x** speedup over the reference
- **Overall** achieved **~3x** speedup

Appendix


Experimental Results (Intel i7-8700K @ 4.4GHz with GCC)

Performance MeasureFull
(Measured with rdtsc; Peak Performance ethSIFT: 3.21)



Experimental Results (Intel i7-8700K @ 4.4GHz with CLANG)

Performance MeasureFull
(Measured with rdtsc; Peak Performance ethSIFT: 3.24)



STEP 1: Scale-space Extrema Detection

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

- Layers in Gaussian Pyramid


$$L(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$$

- Layers in Difference of Gaussians

$$D(x, y, \sigma) = L(x, y, k\sigma) - L(x, y, \sigma)$$

Constant multiplicative factor k


After each octave, image is downsampled by a factor of 2



STEP 1: Scale-space Extrema Detection

Find local extrema in DoG pyramid by comparing a pixel to its 8 neighbour pixels in the same scale, as well as to the 9 pixels in each of the adjacent scales.

It is only selected if it is either larger or smaller than all of these other pixel values.



STEP 2: Keypoint Localization

Refinement of potential keypoint candidates:

- Determine interpolated location of extrema, using Taylor Expansion of scale-space function $D(x, y, \sigma)$

$$x = (x, y, \sigma)^T \quad D(x) = D + \frac{\partial D^T}{\partial x} x + \frac{1}{2} x^T \frac{\partial^2 D}{\partial x^2} x \quad \hat{x} = -\frac{\partial^2 D^{-1}}{\partial x} \frac{\partial D}{\partial x}$$

- Where \hat{x} is the derived location of the extremum. Remove a keypoint if $|D(\hat{x})| < "threshold"$
- For stability, edge responses get eliminated. This is done by first computing the Hessian H :

$$H = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$$

- And use the Hessian to check if ratio of principal curvatures is below some curvature threshold r

$$\frac{Tr(H)^2}{Det(H)} < \frac{(r+1)^2}{r}$$

STEP 3: Orientation Assignment

Next, Orientation is assigned to each keypoint to achieve invariance to image rotation.

- Calculation of gradient magnitude $m(x, y)$, and orientation $\Theta(x, y)$, using pixel differences:

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$
$$\Theta(x, y) = \tan^{-1} \frac{L(x, y+1) - L(x, y-1)}{L(x+1, y) - L(x-1, y)}$$

- An orientation histogram is formed from gradient orientations of sample points within a region around the keypoint. It has 36 bins covering the 360° range of orientations.
- Each sample added to histogram is weighted by its gradient magnitude.
- Peaks in the orientation histogram correspond to dominant directions of local gradients.

STEP 4: Extraction of Keypoint Descriptor

Last step is to create a descriptor for the local image region:

- We consider a 16x16 pixel patch around the keypoint
- The patch is subdivided in 16 blocks of size 4x4, for each of these blocks an 8 bin orientation histogram is created
- This gives us a total of 128 bin values. These values are represented as a vector to form the keypoint descriptor





Figure: Shows a 2x2 descriptor array computed from an 8x8 set of samples



References

- Distinctive Image Features from Scale-Invariant Keypoints [Lowe, 2004]
- Guohui Wang, ezSIFT: an easy-to-use standalone SIFT library, 2013